We picked seed 15, Compute Costs: which stocks are exposed to GPU prices, with Ornn as the industry partner. This one was an easy choice for us. This past summer I worked at a firm focused on hedging for data centers using GPU prices, so I've already seen how much compute costs matter and how hard they are to hedge. We also know Harrison Keen, who worked at Ornn this past summer, so we have a direct connection to the company and a good sense of what they'd want out of the report. It felt like the project where we could actually build something useful for a real client instead of just a class assignment.

We set up the repo with a fake Ornn data file that has the same columns as the real data, plus code that pulls in the stock data, so the whole pipeline runs for anyone who clones it. We also downloaded real GPU price data to our own computers through our Ornn free trial. That data stays local and is not in the repo, per the data rule. You can see the fake Ornn file and what the code produces. The main things we're still figuring out are where to get the stock data and how much GPU history we need.

Questions for you

1. What's the best way to get the stock data? We're using yfinance right now, but we weren't sure if there's a better or more reliable source we should use.
2. How far back should the Ornn GPU data go? The seed says about two years, but we wanted to check whether that's enough, or whether we should push for more history.
3. What's the best way to review the outputs Claude gives us when coding things for GitHub? We want to make sure we actually understand and check what it writes before we push it, especially since every teammate needs to be able to explain each step.